



Prompt Caching

Search Engine Optimization.

"Wireless Bluetooth Headphones" -> LLM to generate SEO optimized keywords, and description.




Behavior	GPT-5.6 and later	GPT-5.5 and GPT-5.5 Pro	Other earlier models
Implicit breakpoints	At the end of the latest eligible user or tool message.	Spaced at regular 2,048-token intervals.	Spaced at regular, model-dependent intervals.
Explicit breakpoints	Supported	Not supported	Not supported
Minimum cacheable prefix	1,024 visible input tokens	2,048 visible input tokens; some models may cache shorter prefixes	2,048 visible input tokens; some models may cache shorter prefixes



req1:
sys_prompt = You're an assistant -> 1200 token, cost is ther

req2:
sys_prompt = ... -> not costed.



Excludes hidden tokens and rounds
down to a multiple of 128

1100 tokens $\rightarrow$ 1024 tokens cached,
76 " are not cached,

1300 tokens $\rightarrow$ 1024 cached,
1280
+ 128 ✓
+ 128 ✓
+ 20 not cached,

1500 - 1408 cached
- 92 not cached.

Text tokens

Per 1M tokens · Batch API price



Input

\$0.05

Cached input

\$0.005

Output

\$0.40

Text tokens

Per 1M tokens · Batch API price



Input

\$0.025

Cached input

\$0.0025

Output


\$0.20

Giant

High throughput Comp.

→ initialized
into queue

Knowledge Distillation



GEN AI Engineer

Checkout requests started failing and API latency has increased suddenly.


1. Restarting some random services.
2. Clearing Caches.
3. Try diff dashboards.
4. Grep logs in 12 different servers, for finding out info on Request failing.
5. Sit in corner and starting chanting mantra.
6. Search slack messages.
7. Apply patch which worked in past.

" When latency is high, restart Service A, B and then G. Then inspect DB, then scale production servers"

Scenario 2:

Senior engineer on day 1:

1. This pattern means DB connection pool Exhaustion.
2. Do not restart anything until you've verified exhaustion.
3. Check these 3 metrics and start service B.



Steps Distillation.

We ask the strong model to generate the solution and the sequence of steps it followed to arrive at the solution.

we also ask the model to narrate it's chain of understanding, the sequence of steps it followed.

book a flight ticket.

1. booking.com
2. list accessibility tree
3. look for signin ...
- 4.

